

# Helper\_function\_implementations

January 26, 2025

## 1 KNN.py

```
[ ]: import numpy as np
import pdb

class KNN(object):

    def __init__(self):
        pass

    def train(self, X, y):
        """
        Inputs:
        - X is a numpy array of size (num_examples, D)
        - y is a numpy array of size (num_examples, )
        """
        self.X_train = X
        self.y_train = y

    def compute_distances(self, X, norm=None):
        """
        Compute the distance between each test point in X and each training point
        in self.X_train.

        Inputs:
        - X: A numpy array of shape (num_test, D) containing test data.
        - norm: the function with which the norm is taken.

        Returns:
        - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
            is the Euclidean distance between the ith test point and the jth training
            point.
        """
        if norm is None:
            norm = lambda x: np.sqrt(np.sum(x**2))
            #norm = 2
```

```

num_test = X.shape[0]
num_train = self.X_train.shape[0]
dists = np.zeros((num_test, num_train))
for i in np.arange(num_test):

    for j in np.arange(num_train):
        # ===== #
        # YOUR CODE HERE:
        # Compute the distance between the ith test point and the jth
        # training point using norm(), and store the result in dists[i, j].

        # ===== #
        distance = X[i, :] - self.X_train[j, :]
        dists[i, j] = norm(distance)

        # ===== #
        # END YOUR CODE HERE
        # ===== #

    return dists

def compute_L2_distances_vectorized(self, X):
    """
    Compute the distance between each test point in X and each training point
    in self.X_train WITHOUT using any for loops.

    Inputs:
    - X: A numpy array of shape (num_test, D) containing test data.

    Returns:
    - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
      is the Euclidean distance between the ith test point and the jth training
      point.
    """
    num_test = X.shape[0]
    num_train = self.X_train.shape[0]
    dists = np.zeros((num_test, num_train))

    # ===== #
    # YOUR CODE HERE:
    # Compute the L2 distance between the ith test point and the jth
    # training point and store the result in dists[i, j]. You may
    # NOT use a for loop (or list comprehension). You may only use
    # numpy operations.
    #
    # HINT: use broadcasting. If you have a shape (N,1) array and

```

```

# a shape (M,) array, adding them together produces a shape (N, M)
# array.
# ===== #
dists = np.sqrt(np.sum(X**2, axis = 1).reshape((num_test, 1)) + np.sum(self.
↪X_train**2, axis=1) - 2*np.dot(X,self.X_train.T))

# ===== #
# END YOUR CODE HERE
# ===== #

return dists

def predict_labels(self, dists, k=1):
    """
    Given a matrix of distances between test points and training points,
    predict a label for each test point.

    Inputs:
    - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
      gives the distance between the ith test point and the jth training point.

    Returns:
    - y: A numpy array of shape (num_test,) containing predicted labels for the
      test data, where y[i] is the predicted label for the test point X[i].
    """
    num_test = dists.shape[0]
    y_pred = np.zeros(num_test)
    for i in np.arange(num_test):
        # A list of length k storing the labels of the k nearest neighbors to
        # the ith test point.
        closest_y = []
        # ===== #
        # YOUR CODE HERE:
        # Use the distances to calculate and then store the labels of
        # the k-nearest neighbors to the ith test point. The function
        # numpy.argsort may be useful.
        #
        # After doing this, find the most common label of the k-nearest
        # neighbors. Store the predicted label of the ith training example
        # as y_pred[i]. Break ties by choosing the smaller label.
        # ===== #

        sorted = np.argsort(dists[i,:])
        closest_y = self.y_train[sorted[:k]]
        y_pred[i] = np.argmax(np.bincount(closest_y))

```

```

# ===== #
# END YOUR CODE HERE
# ===== #

return y_pred

```

## 2 Softmax.py

```

[ ]: import numpy as np

class Softmax(object):

    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)

    def init_weights(self, dims):
        """
        Initializes the weight matrix of the Softmax classifier.
        Note that it has shape (C, D) where C is the number of
        classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims) * 0.0001

    def loss(self, X, y):
        """
        Calculates the softmax loss.

        Inputs have dimension D, there are C classes, and we operate on minibatches
        of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i] = c means
            that X[i] has label c, where 0 <= c < C.

        Returns a tuple of:
        - loss as single float
        """
        # Initialize the loss to zero.
        loss = 0.0

        # ===== #
        # YOUR CODE HERE:
        # Calculate the normalized softmax loss. Store it as the variable loss.

```

```

# (That is, calculate the sum of the losses of all the training
# set margins, and then normalize the loss by the number of
# training examples.)
# ===== #
N,D = X.shape
all_scores = np.dot(X, self.W.T)
for i in range(N):
    score_i = all_scores[i,:]
    score_i -= np.max(score_i)
    current_score = score_i[y[i]]
    loss += current_score - np.log(np.sum(np.exp(score_i)))
loss /= -N
# ===== #
# END YOUR CODE HERE
# ===== #

return loss

def loss_and_grad(self, X, y):
    """
    Same as self.loss(X, y), except that it also returns the gradient.

    Output: grad -- a matrix of the same dimensions as W containing
        the gradient of the loss with respect to W.
    """

    # Initialize the loss and gradient to zero.
    loss = 0.0
    grad = np.zeros_like(self.W)

    # ===== #
    # YOUR CODE HERE:
    # Calculate the softmax loss and the gradient. Store the gradient
    # as the variable grad.
    # ===== #
    N = X.shape[0]
    M = self.W.shape[0]
    all_scores = np.dot(X, self.W.T)
    for i in range(N):
        score_i = all_scores[i,:]
        score_i -= np.max(score_i)
        grad[y[i]] += X[i]
        for m in range(M):
            grad[m] -= np.exp(score_i[m])/np.sum(np.exp(score_i))*X[i]
    loss = self.loss(X,y)
    grad /= -N
    # ===== #

```

```

# END YOUR CODE HERE
# ===== #

return loss, grad

def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
    """
    sample a few random elements and only return numerical
    in these dimensions.
    """

    for i in np.arange(num_checks):
        ix = tuple([np.random.randint(m) for m in self.W.shape])

        oldval = self.W[ix]
        self.W[ix] = oldval + h # increment by h
        fxph = self.loss(X, y)
        self.W[ix] = oldval - h # decrement by h
        fxmh = self.loss(X,y) # evaluate f(x - h)
        self.W[ix] = oldval # reset

        grad_numerical = (fxph - fxmh) / (2 * h)
        grad_analytic = your_grad[ix]
        rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) +
↪abs(grad_analytic))
        print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical,
↪grad_analytic, rel_error))

def fast_loss_and_grad(self, X, y):
    """
    A vectorized implementation of loss_and_grad. It shares the same
    inputs and ouputs as loss_and_grad.
    """

    loss = 0.0
    grad = np.zeros(self.W.shape) # initialize the gradient as zero

    # ===== #
    # YOUR CODE HERE:
    # Calculate the softmax loss and gradient WITHOUT any for loops.
    # ===== #
    N = X.shape[0] # Number of samples
    C = self.W.shape[0] # Number of classes

    # Compute scores
    scores = np.dot(X, self.W.T) # Shape: (N, C)
    scores -= np.max(scores, axis=1, keepdims=True) # For numerical stability

```

```

# Compute softmax probabilities
exp_scores = np.exp(scores) # Shape: (N, C)
probabilities = exp_scores / np.sum(exp_scores, axis=1, keepdims=True) #
↳ Shape: (N, C)

# Compute loss
correct_log_probs = -np.log(probabilities[np.arange(N), y]) # Shape: (N,)
loss = np.sum(correct_log_probs) / N # Scalar loss

# Compute gradient
probabilities[np.arange(N), y] -= 1 # Subtract 1 for correct class
↳ probabilities
grad = np.dot(probabilities.T, X) / N # Shape: (C, D)
# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
          batch_size=200, verbose=False):
    """
    Train this linear classifier using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) containing training data; there are N
        training samples each of dimension D.
    - y: A numpy array of shape (N,) containing training labels; y[i] = c
        means that X[i] has label 0 ≤ c < C for C classes.
    - learning_rate: (float) learning rate for optimization.
    - num_iters: (integer) number of steps to take when optimizing
    - batch_size: (integer) number of training examples to use at each step.
    - verbose: (boolean) If true, print progress during optimization.

    Outputs:
    A list containing the value of the loss function at each training iteration.
    """
    num_train, dim = X.shape
    num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is
    ↳ number of classes

    self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes
    ↳ the weights of self.W

    # Run stochastic gradient descent to optimize W
    loss_history = []

```

```

for it in np.arange(num_iters):
    X_batch = None
    y_batch = None

    # ===== #
    # YOUR CODE HERE:
    # Sample batch_size elements from the training data for use in
    # gradient descent. After sampling,
    # - X_batch should have shape: (batch_size, dim)
    # - y_batch should have shape: (batch_size,)
    # The indices should be randomly generated to reduce correlations
    # in the dataset. Use np.random.choice. It's okay to sample with
    # replacement.
    # ===== #
    indexes = np.random.choice(num_train, batch_size)
    X_batch = X[indexes]
    y_batch = y[indexes]
    # ===== #
    # END YOUR CODE HERE
    # ===== #
    # evaluate loss and gradient
    loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
    loss_history.append(loss)

    # ===== #
    # YOUR CODE HERE:
    # Update the parameters, self.W, with a gradient step
    # ===== #
    self.W -= learning_rate*grad

    # ===== #
    # END YOUR CODE HERE
    # ===== #
    if verbose and it % 100 == 0:
        print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

return loss_history

def predict(self, X):
    """
    Inputs:
    - X: N x D array of training data. Each row is a D-dimensional point.

    Returns:
    - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
      array of length N, and each element is an integer giving the predicted

```



```

class.
"""
y_pred = np.zeros(X.shape[1])
# ===== #
# YOUR CODE HERE:
#   Predict the labels given the training data.
# ===== #
all_scores = np.dot(X, self.W.T)
all_scores = (all_scores.T - np.max(all_scores, axis = 1)).T
probabilities = np.exp(all_scores)/np.sum(np.exp(all_scores), axis = 1,
↳keepdims = True)
y_pred = np.argmax(probabilities, axis = 1) #y_pred should be a 1 dim array↳
↳of length N
# ===== #
# END YOUR CODE HERE
# ===== #

return y_pred

```