

# HW3\_Q4\_Q5\_Helper\_Files

February 7, 2025

## 1 Question 4 - 2 Layer Neural Network

### 1.1 data\_utils.py

Complete the two-layer neural network Jupyter notebook. Print out the entire notebook and relevant code and submit it as a pdf to gradescope. Download the CIFAR-10 dataset, as you did in HW #2.

```
[4]: from __future__ import print_function

from six.moves import cPickle as pickle
import numpy as np
import os
from matplotlib.pyplot import imread
import platform

def load_pickle(f):
    version = platform.python_version_tuple()
    if version[0] == '2':
        return pickle.load(f)
    elif version[0] == '3':
        return pickle.load(f, encoding='latin1')
    raise ValueError("invalid python version: {}".format(version))

def load_CIFAR_batch(filename):
    """ load single batch of cifar """
    with open(filename, 'rb') as f:
        datadict = load_pickle(f)
        X = datadict['data']
        Y = datadict['labels']
        X = X.reshape(10000, 3, 32, 32).transpose(0,2,3,1).astype("float")
        Y = np.array(Y)
        return X, Y

# def load_CIFAR10(ROOT):
#     """ load all of cifar """
#     xs = []
#     ys = []
```

```

#     for b in range(1,6):
#         f = os.path.join(ROOT, 'data_batch_%d' % (b, ))
#         X, Y = load_CIFAR_batch(f)
#         xs.append(X)
#         ys.append(Y)
#     Xtr = np.concatenate(xs)
#     Ytr = np.concatenate(ys)
#     del X, Y
#     Xte, Yte = load_CIFAR_batch(os.path.join(ROOT, 'test_batch'))
#     return Xtr, Ytr, Xte, Yte

def load_CIFAR10(ROOT):
    """Load all of CIFAR-10 using absolute paths."""
    xs = []
    ys = []
    """ NOTE FOR THE GRADERS: I had something going on with my join ROOT
    ↪function
        so I decided to simply manually join the file directory with a similar
    ↪for loop
        because I kept getting the same error despite having the correct
    ↪directory
        You can see I tested to see if the directory is present in the normal
    ↪code"""
    for b in range(1, 6):
        f = f"/Users/ctang/Desktop/ECE_C147/HW3/cifar-10-batches-py/
    ↪data_batch_{b}"
        if not os.path.exists(f):
            raise FileNotFoundError(f"File not found: {f}")
        X, Y = load_CIFAR_batch(f)
        xs.append(X)
        ys.append(Y)

    Xtr = np.concatenate(xs)
    Ytr = np.concatenate(ys)
    del X, Y

    test_file = "/Users/ctang/Desktop/ECE_C147/HW3/cifar-10-batches-py/
    ↪test_batch"
    if not os.path.exists(test_file):
        raise FileNotFoundError(f"File not found: {test_file}")

    Xte, Yte = load_CIFAR_batch(test_file)
    return Xtr, Ytr, Xte, Yte

def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000,

```

```

        subtract_mean=True):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for classifiers. These are the same steps as we used for the SVM, but
    condensed to a single function.
    """
    # Load the raw CIFAR-10 data
    cifar10_dir = '/Users/ctang/Desktop/ECE_C147/HW3/cifar-10-batches-py/'
    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # Subsample the data
    mask = list(range(num_training, num_training + num_validation))
    X_val = X_train[mask]
    y_val = y_train[mask]
    mask = list(range(num_training))
    X_train = X_train[mask]
    y_train = y_train[mask]
    mask = list(range(num_test))
    X_test = X_test[mask]
    y_test = y_test[mask]

    # Normalize the data: subtract the mean image
    if subtract_mean:
        mean_image = np.mean(X_train, axis=0)
        X_train -= mean_image
        X_val -= mean_image
        X_test -= mean_image

    # Transpose so that channels come first
    X_train = X_train.transpose(0, 3, 1, 2).copy()
    X_val = X_val.transpose(0, 3, 1, 2).copy()
    X_test = X_test.transpose(0, 3, 1, 2).copy()

    # Package data into a dictionary
    return {
        'X_train': X_train, 'y_train': y_train,
        'X_val': X_val, 'y_val': y_val,
        'X_test': X_test, 'y_test': y_test,
    }

def load_tiny_imagenet(path, dtype=np.float32, subtract_mean=True):
    """
    Load TinyImageNet. Each of TinyImageNet-100-A, TinyImageNet-100-B, and
    TinyImageNet-200 have the same directory structure, so this can be used
    to load any of them.

```

*Inputs:*

- *path*: String giving path to the directory to load.
- *dtype*: numpy datatype used to load the data.
- *subtract\_mean*: Whether to subtract the mean training image.

*Returns: A dictionary with the following entries:*

- *class\_names*: A list where *class\_names[i]* is a list of strings giving the WordNet names for class *i* in the loaded dataset.
- *X\_train*: (*N\_tr*, 3, 64, 64) array of training images
- *y\_train*: (*N\_tr*,) array of training labels
- *X\_val*: (*N\_val*, 3, 64, 64) array of validation images
- *y\_val*: (*N\_val*,) array of validation labels
- *X\_test*: (*N\_test*, 3, 64, 64) array of testing images.
- *y\_test*: (*N\_test*,) array of test labels; if test labels are not available (such as in student code) then *y\_test* will be None.
- *mean\_image*: (3, 64, 64) array giving mean training image

*# First load wnids*

```
with open(os.path.join(path, 'wnids.txt'), 'r') as f:
    wnids = [x.strip() for x in f]
```

*# Map wnids to integer labels*

```
wnid_to_label = {wnid: i for i, wnid in enumerate(wnids)}
```

*# Use words.txt to get names for each class*

```
with open(os.path.join(path, 'words.txt'), 'r') as f:
    wnid_to_words = dict(line.split('\t') for line in f)
    for wnid, words in wnid_to_words.iteritems():
        wnid_to_words[wnid] = [w.strip() for w in words.split(',')]
class_names = [wnid_to_words[wnid] for wnid in wnids]
```

*# Next load training data.*

```
X_train = []
```

```
y_train = []
```

```
for i, wnid in enumerate(wnids):
```

```
    if (i + 1) % 20 == 0:
```

```
        print('loading training data for synset %d / %d' % (i + 1,
↪len(wnids)))
```

*# To figure out the filenames we need to open the boxes file*

```
boxes_file = os.path.join(path, 'train', wnid, '%s_boxes.txt' % wnid)
```

```
with open(boxes_file, 'r') as f:
```

```
    filenames = [x.split('\t')[0] for x in f]
```

```
num_images = len(filenames)
```

```
X_train_block = np.zeros((num_images, 3, 64, 64), dtype=dtype)
```

```
y_train_block = wnid_to_label[wnid] * np.ones(num_images, dtype=np.int64)
```

```
for j, img_file in enumerate(filenames):
```

```

img_file = os.path.join(path, 'train', wnid, 'images', img_file)
img = imread(img_file)
if img.ndim == 2:
    ## grayscale file
    img.shape = (64, 64, 1)
X_train_block[j] = img.transpose(2, 0, 1)
X_train.append(X_train_block)
y_train.append(y_train_block)

# We need to concatenate all training data
X_train = np.concatenate(X_train, axis=0)
y_train = np.concatenate(y_train, axis=0)

# Next load validation data
with open(os.path.join(path, 'val', 'val_annotations.txt'), 'r') as f:
    img_files = []
    val_wnids = []
    for line in f:
        img_file, wnid = line.split('\t')[:2]
        img_files.append(img_file)
        val_wnids.append(wnid)
    num_val = len(img_files)
    y_val = np.array([wnid_to_label[wnid] for wnid in val_wnids])
    X_val = np.zeros((num_val, 3, 64, 64), dtype=dtype)
    for i, img_file in enumerate(img_files):
        img_file = os.path.join(path, 'val', 'images', img_file)
        img = imread(img_file)
        if img.ndim == 2:
            img.shape = (64, 64, 1)
        X_val[i] = img.transpose(2, 0, 1)

# Next load test images
# Students won't have test labels, so we need to iterate over files in the
# images directory.
img_files = os.listdir(os.path.join(path, 'test', 'images'))
X_test = np.zeros((len(img_files), 3, 64, 64), dtype=dtype)
for i, img_file in enumerate(img_files):
    img_file = os.path.join(path, 'test', 'images', img_file)
    img = imread(img_file)
    if img.ndim == 2:
        img.shape = (64, 64, 1)
    X_test[i] = img.transpose(2, 0, 1)

y_test = None
y_test_file = os.path.join(path, 'test', 'test_annotations.txt')
if os.path.isfile(y_test_file):
    with open(y_test_file, 'r') as f:

```

```

        img_file_to_wnid = {}
        for line in f:
            line = line.split('\t')
            img_file_to_wnid[line[0]] = line[1]
        y_test = [wnid_to_label[img_file_to_wnid[img_file]] for img_file in
img_files]
        y_test = np.array(y_test)

    mean_image = X_train.mean(axis=0)
    if subtract_mean:
        X_train -= mean_image[None]
        X_val -= mean_image[None]
        X_test -= mean_image[None]

    return {
        'class_names': class_names,
        'X_train': X_train,
        'y_train': y_train,
        'X_val': X_val,
        'y_val': y_val,
        'X_test': X_test,
        'y_test': y_test,
        'class_names': class_names,
        'mean_image': mean_image,
    }

def load_models(models_dir):
    """
    Load saved models from disk. This will attempt to unpickle all files in a
    directory; any files that give errors on unpickling (such as README.txt)
will
    be skipped.

    Inputs:
    - models_dir: String giving the path to a directory containing model files.
    Each model file is a pickled dictionary with a 'model' field.

    Returns:
    A dictionary mapping model file names to models.
    """
    models = {}
    for model_file in os.listdir(models_dir):
        with open(os.path.join(models_dir, model_file), 'rb') as f:
            try:
                models[model_file] = load_pickle(f)['model']
            except pickle.UnpicklingError:

```

```

        continue
    return models

```

## 1.2 neural\_net.py

```

[7]: import numpy as np
import matplotlib.pyplot as plt

class TwoLayerNet(object):
    """
    A two-layer fully-connected neural network. The net has an input dimension  $D$  of
     $D$ , a hidden layer dimension of  $H$ , and performs classification over  $C$ 
    classes.
    We train the network with a softmax loss function and L2 regularization on the
    weight matrices. The network uses a ReLU nonlinearity after the first fully
    connected layer.

    In other words, the network has the following architecture:

    input - fully connected layer - ReLU - fully connected layer - softmax

    The outputs of the second fully-connected layer are the scores for each
    class.
    """

    def __init__(self, input_size, hidden_size, output_size, std=1e-4):
        """
        Initialize the model. Weights are initialized to small random values and
        biases are initialized to zero. Weights and biases are stored in the
        variable self.params, which is a dictionary with the following keys:

        W1: First layer weights; has shape (H, D)
        b1: First layer biases; has shape (H,)
        W2: Second layer weights; has shape (C, H)
        b2: Second layer biases; has shape (C,)

        Inputs:
        - input_size: The dimension  $D$  of the input data.
        - hidden_size: The number of neurons  $H$  in the hidden layer.
        - output_size: The number of classes  $C$ .
        """
        # self.params = {}
        # self.params['W1'] = std * np.random.randn(hidden_size, input_size)
        # self.params['b1'] = np.zeros(hidden_size)

```

```

# self.params['W2'] = std * np.random.randn(output_size, hidden_size)
# self.params['b2'] = np.zeros(output_size)
self.params = {}
self.params['W1'] = std * np.random.randn(hidden_size, input_size)
self.params['b1'] = np.zeros(hidden_size)
self.params['W2'] = std * np.random.randn(output_size, hidden_size)
self.params['b2'] = np.zeros(output_size)

def loss(self, X, y=None, reg=0.0):
    """
    Compute the loss and gradients for a two layer fully connected neural
    network.

    Inputs:
    - X: Input data of shape (N, D). Each X[i] is a training sample.
    - y: Vector of training labels. y[i] is the label for X[i], and each
    ↪ y[i] is
        an integer in the range 0 ≤ y[i] < C. This parameter is optional; if it
        is not passed then we only return scores, and if it is passed then we
        instead return the loss and gradients.
    - reg: Regularization strength.

    Returns:
    If y is None, return a matrix scores of shape (N, C) where scores[i, c]
    ↪ is
        the score for class c on input X[i].

    If y is not None, instead return a tuple of:
    - loss: Loss (data loss and regularization loss) for this batch of
    ↪ training
        samples.
    - grads: Dictionary mapping parameter names to gradients of those
    ↪ parameters
        with respect to the loss function; has the same keys as self.params.
    """
    # Unpack variables from the params dictionary
    W1, b1 = self.params['W1'], self.params['b1']
    W2, b2 = self.params['W2'], self.params['b2']
    N, D = X.shape

    # Compute the forward pass
    scores = None
    h1 = np.dot(X, W1.T) + b1
    h1[h1 ≤ 0] = 0
    h2 = np.dot(h1, W2.T) + b2
    # ===== #
    # YOUR CODE HERE:

```



```

# Calculate the output scores of the neural network. The result
# should be (N, C). As stated in the description for this class,
# there should not be a ReLU layer after the second FC layer.
# The output of the second FC layer is the output scores. Do not
# use a for loop in your implementation.
# ===== #
real_scores = h2
scores = h2 - np.max(h2, axis=1, keepdims=True)
exp_scores = np.exp(scores)
probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)
#pass

# ===== #
# END YOUR CODE HERE
# ===== #

# If the targets are not given then jump out, we're done
if y is None:
    #return scores
    return real_scores

# Compute the loss
loss = None

# ===== #
# YOUR CODE HERE:
# Calculate the loss of the neural network. This includes the
# softmax loss and the L2 regularization for W1 and W2. Store the
# total loss in the variable loss. Multiply the regularization
# loss by 0.5 (in addition to the factor reg).
# ===== #
# scores is num_examples by num_classes
num_examples = X.shape[0]
correct_logprobs = -np.log(probs[range(num_examples), y])
data_loss = np.sum(correct_logprobs) / num_examples
reg_loss = 0.5*reg*(np.sum(W1 * W1) + np.sum(W2 * W2))
loss = data_loss + reg_loss
#pass
# ===== #
# END YOUR CODE HERE
# ===== #

grads = {}

# ===== #
# YOUR CODE HERE:

```

```

# Implement the backward pass. Compute the derivatives of the
# weights and the biases. Store the results in the grads
# dictionary. e.g., grads['W1'] should store the gradient for
# W1, and be of the same size as W1.
# ===== #
dscores = probs
dscores[range(N), y] -= 1
dscores /= N
grads['W2'] = np.dot(dscores.T, h1) + reg*W2
grads['b2'] = np.sum(dscores, axis=0)
dh1 = np.dot(dscores, W2)
dh1[h1<=0]=0
grads['W1'] = np.dot(dh1.T, X) + reg * W1
grads['b1'] = np.sum(dh1, axis=0)

#pass

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grads

def train(self, X, y, X_val, y_val,
          learning_rate=1e-3, learning_rate_decay=0.95,
          reg=1e-5, num_iters=100,
          batch_size=200, verbose=False):
    """
    Train this neural network using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) giving training data.
    - y: A numpy array of shape (N,) giving training labels; y[i] = c means
    ↪ that X[i] has label c, where 0 ≤ c < C.
    - X_val: A numpy array of shape (N_val, D) giving validation data.
    - y_val: A numpy array of shape (N_val,) giving validation labels.
    - learning_rate: Scalar giving learning rate for optimization.
    - learning_rate_decay: Scalar giving factor used to decay the learning
    ↪ rate after each epoch.
    - reg: Scalar giving regularization strength.
    - num_iters: Number of steps to take when optimizing.
    - batch_size: Number of training examples to use per step.
    - verbose: boolean; if true print progress during optimization.
    """
    num_train = X.shape[0]

```

```

iterations_per_epoch = max(num_train / batch_size, 1)

# Use SGD to optimize the parameters in self.model
loss_history = []
train_acc_history = []
val_acc_history = []

for it in np.arange(num_iters):
    X_batch = None
    y_batch = None

    # ===== #
    # YOUR CODE HERE:
    # Create a minibatch by sampling batch_size samples randomly.
    # ===== #
    #pass
    indexes = np.random.choice(num_train, batch_size)
    X_batch = X[indexes]
    y_batch = y[indexes]

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    # Compute loss and gradients using the current minibatch
    loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
    loss_history.append(loss)

    # ===== #
    # YOUR CODE HERE:
    # Perform a gradient descent step using the minibatch to update
    # all parameters (i.e., W1, W2, b1, and b2).
    # ===== #

    self.params['W1'] -= learning_rate*grads['W1']
    self.params['W2'] -= learning_rate*grads['W2']
    self.params['b1'] -= learning_rate*grads['b1']
    self.params['b2'] -= learning_rate*grads['b2']
    #pass

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    if verbose and it % 100 == 0:
        print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

```

```

        # Every epoch, check train and val accuracy and decay learning rate.
        if it % iterations_per_epoch == 0:
            # Check accuracy
            train_acc = (self.predict(X_batch) == y_batch).mean()
            val_acc = (self.predict(X_val) == y_val).mean()
            train_acc_history.append(train_acc)
            val_acc_history.append(val_acc)

            # Decay learning rate
            learning_rate *= learning_rate_decay

    return {
        'loss_history': loss_history,
        'train_acc_history': train_acc_history,
        'val_acc_history': val_acc_history,
    }

def predict(self, X):
    """
    Use the trained weights of this two-layer network to predict labels for
    data points. For each data point we predict scores for each of the C
    classes, and assign each data point to the class with the highest score.

    Inputs:
    - X: A numpy array of shape (N, D) giving N D-dimensional data points to
        classify.

    Returns:
    - y_pred: A numpy array of shape (N,) giving predicted labels for each
    ↪ of
        the elements of X. For all i, y_pred[i] = c means that X[i] is
    ↪ predicted
        to have class c, where 0 ≤ c < C.
    """
    y_pred = None

    # ===== #
    # YOUR CODE HERE:
    # Predict the class given the input data.
    # ===== #
    #pass
    predicted_output = np.dot(np.maximum(0, np.dot(X, self.params['W1']).T) +
    ↪ self.params['b1']), self.params['W2']).T + self.params['b2']
    y_pred = np.argmax(predicted_output, axis=1)

```

```

# ===== #
# END YOUR CODE HERE
# ===== #

return y_pred

```

## 2 Question 5 FC Nets

Complete the FC Net Jupyter notebook. Print out the entire notebook and relevant code and submit it as a pdf to gradescope

### 2.1 layers.py

```

[14]: import numpy as np
import pdb

def affine_forward(x, w, b):
    """
    Computes the forward pass for an affine (fully-connected) layer.

    The input x has shape (N, d_1, ..., d_k) and contains a minibatch of N
    examples, where each example x[i] has shape (d_1, ..., d_k). We will
    reshape each input into a vector of dimension D = d_1 * ... * d_k, and
    then transform it to an output vector of dimension M.

    Inputs:
    - x: A numpy array containing input data, of shape (N, d_1, ..., d_k)
    - w: A numpy array of weights, of shape (D, M)
    - b: A numpy array of biases, of shape (M,)

    Returns a tuple of:
    - out: output, of shape (N, M)
    - cache: (x, w, b)
    """

    # ===== #
    # YOUR CODE HERE:
    #   Calculate the output of the forward pass. Notice the dimensions
    #   of w are D x M, which is the transpose of what we did in earlier
    #   assignments.
    # ===== #
    X = x.reshape((x.shape[0], -1))
    out = np.dot(X, w) + b
    #pass

    # ===== #
    # END YOUR CODE HERE

```

```

# ===== #

cache = (x, w, b)
return out, cache

def affine_backward(dout, cache):
    """
    Computes the backward pass for an affine layer.

    Inputs:
    - dout: Upstream derivative, of shape (N, M)
    - cache: Tuple of:
    - x: Input data, of shape (N, d_1, ... d_k)
    - w: Weights, of shape (D, M)

    Returns a tuple of:
    - dx: Gradient with respect to x, of shape (N, d1, ..., d_k)
    - dw: Gradient with respect to w, of shape (D, M)
    - db: Gradient with respect to b, of shape (M,)
    """
    x, w, b = cache
    dx, dw, db = None, None, None

    # ===== #
    # YOUR CODE HERE:
    # Calculate the gradients for the backward pass.
    # ===== #

    # dout is N x M
    # dx should be N x d1 x ... x dk; it relates to dout through multiplication
    # with w, which is D x M
    # dw should be D x M; it relates to dout through multiplication with x,
    # which is N x D after reshaping
    # db should be M; it is just the sum over dout examples
    X = x.reshape((x.shape[0], -1))
    db = np.sum(dout, axis = 0)
    dw = np.dot(X.T, dout)
    dx = np.dot(dout, w.T).reshape(x.shape)
    #pass

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return dx, dw, db

```

```

def relu_forward(x):
    """
    Computes the forward pass for a layer of rectified linear units (ReLU).

    Input:
    - x: Inputs, of any shape

    Returns a tuple of:
    - out: Output, of the same shape as x
    - cache: x
    """
    # ===== #
    # YOUR CODE HERE:
    # Implement the ReLU forward pass.
    # ===== #
    out = np.maximum(0,x)
    #pass
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    cache = x
    return out, cache

def relu_backward(dout, cache):
    """
    Computes the backward pass for a layer of rectified linear units (ReLU).

    Input:
    - dout: Upstream derivatives, of any shape
    - cache: Input x, of same shape as dout

    Returns:
    - dx: Gradient with respect to x
    """
    x = cache

    # ===== #
    # YOUR CODE HERE:
    # Implement the ReLU backward pass
    # ===== #

    # ReLU directs linearly to those > 0
    #pass
    dx = dout*(x>0)

```

```

# ===== #
# END YOUR CODE HERE
# ===== #

return dx

def softmax_loss(x, y):
    """
    Computes the loss and gradient for softmax classification.

    Inputs:
    - x: Input data, of shape (N, C) where x[i, j] is the score for the jth
    ↪ class
      for the ith input.
    - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
      0 ≤ y[i] < C

    Returns a tuple of:
    - loss: Scalar giving the loss
    - dx: Gradient of the loss with respect to x
    """
    probs = np.exp(x - np.max(x, axis=1, keepdims=True))
    probs /= np.sum(probs, axis=1, keepdims=True)
    N = x.shape[0]
    loss = -np.sum(np.log(probs[np.arange(N), y])) / N
    dx = probs.copy()
    dx[np.arange(N), y] -= 1
    dx /= N
    return loss, dx

```

## 2.2 fc\_net.py

```

[ ]: import numpy as np

from .layers import *
from .layer_utils import *

class TwoLayerNet(object):
    """
    A two-layer fully-connected neural network with ReLU nonlinearity and
    softmax loss that uses a modular layer design. We assume an input dimension
    of D, a hidden dimension of H, and perform classification over C classes.

    The architecture should be affine - relu - affine - softmax.

```



*Note that this class does not implement gradient descent; instead, it will interact with a separate Solver object that is responsible for running optimization.*

*The learnable parameters of the model are stored in the dictionary self.params that maps parameter names to numpy arrays.*

```
"""
def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
              dropout=0, weight_scale=1e-3, reg=0.0):
    """
    Initialize a new network.

    Inputs:
    - input_dim: An integer giving the size of the input
    - hidden_dims: An integer giving the size of the hidden layer
    - num_classes: An integer giving the number of classes to classify
    - dropout: Scalar between 0 and 1 giving dropout strength.
    - weight_scale: Scalar giving the standard deviation for random
      initialization of the weights.
    - reg: Scalar giving L2 regularization strength.
    """
    self.params = {}
    self.reg = reg

    # ===== #
    # YOUR CODE HERE:
    #   Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
    #   self.params['W2'], self.params['b1'] and self.params['b2']. The
    #   biases are initialized to zero and the weights are initialized
    #   so that each parameter has mean 0 and standard deviation ↵
    ↪weight_scale.
    #   The dimensions of W1 should be (input_dim, hidden_dim) and the
    #   dimensions of W2 should be (hidden_dims, num_classes)
    # ===== #
    self.params['W1'] = np.random.normal(0, weight_scale, (input_dim, ↵
    ↪hidden_dims))
    self.params['W2'] = np.random.normal(0, weight_scale, (hidden_dims, ↵
    ↪num_classes))
    self.params['b1'] = np.zeros(hidden_dims)
    self.params['b2'] = np.zeros(num_classes)
    #pass

    # ===== #
    # END YOUR CODE HERE
    # ===== #
```

```

def loss(self, X, y=None):
    """
    Compute loss and gradient for a minibatch of data.

    Inputs:
    - X: Array of input data of shape (N, d_1, ..., d_k)
    - y: Array of labels, of shape (N,). y[i] gives the label for X[i].

    Returns:
    If y is None, then run a test-time forward pass of the model and return:
    - scores: Array of shape (N, C) giving classification scores, where
      scores[i, c] is the classification score for X[i] and class c.

    If y is not None, then run a training-time forward and backward pass and
    return a tuple of:
    - loss: Scalar value giving the loss
    - grads: Dictionary with the same keys as self.params, mapping parameter
      names to gradients of the loss with respect to those parameters.
    """
    scores = None

    # ===== #
    # YOUR CODE HERE:
    # Implement the forward pass of the two-layer neural network. Store
    # the class scores as the variable 'scores'. Be sure to use the
    ↪ layers
    # you prior implemented.
    # ===== #
    hidden, cache_hidden = affine_relu_forward(X, self.params['W1'], self.
    ↪ params['b1'])
    scores, cache_scores = affine_forward(hidden, self.params['W2'], self.
    ↪ params['b2'])
    #pass
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    # If y is None then we are in test mode so just return scores
    if y is None:
        return scores

    loss, grads = 0, {}
    # ===== #
    # YOUR CODE HERE:
    # Implement the backward pass of the two-layer neural net. Store
    # the loss as the variable 'loss' and store the gradients in the

```

```

# 'grads' dictionary. For the grads dictionary, grads['W1'] holds
# the gradient for W1, grads['b1'] holds the gradient for b1, etc.
# i.e., grads[k] holds the gradient for self.params[k].
#
# Add L2 regularization, where there is an added cost 0.5*self.reg*W^2
# for each W. Be sure to include the 0.5 multiplying factor to
# match our implementation.
#
# And be sure to use the layers you prior implemented.
# ===== #
loss, dout = softmax_loss(scores,y)
loss += 0.5 * self.reg * (np.sum(self.params['W1']**2) + np.sum(self.
↪params['W2']**2))

dh, dw2, db2 = affine_backward(dout, cache_scores)
dx, dw1, db1 = affine_relu_backward(dh, cache_hidden)

grads['W1'] = dw1 + self.reg * self.params['W1']
grads['b1'] = db1
grads['W2'] = dw2 + self.reg * self.params['W2']
grads['b2'] = db2

#pass

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grads

```

```

class FullyConnectedNet(object):

```

```

    """

```

*A fully-connected neural network with an arbitrary number of hidden layers, ReLU nonlinearities, and a softmax loss function. This will also implement dropout and batch normalization as options. For a network with L layers, the architecture will be*

*{affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax*

*where batch normalization and dropout are optional, and the {...} block is repeated L - 1 times.*

*Similar to the TwoLayerNet above, learnable parameters are stored in the self.params dictionary and will be learned using the Solver class.*

```

    """

```

```

def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
              dropout=0, use_batchnorm=False, reg=0.0,
              weight_scale=1e-2, dtype=np.float32, seed=None):
    """
    Initialize a new FullyConnectedNet.

    Inputs:
    - hidden_dims: A list of integers giving the size of each hidden layer.
    - input_dim: An integer giving the size of the input.
    - num_classes: An integer giving the number of classes to classify.
    - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0,
    then
        the network should not use dropout at all.
    - use_batchnorm: Whether or not the network should use batch
    normalization.
    - reg: Scalar giving L2 regularization strength.
    - weight_scale: Scalar giving the standard deviation for random
        initialization of the weights.
    - dtype: A numpy datatype object; all computations will be performed
    using
        this datatype. float32 is faster but less accurate, so you should use
        float64 for numeric gradient checking.
    - seed: If not None, then pass this random seed to the dropout layers.
    This
        will make the dropout layers deterministic so we can gradient check
    the
        model.
    """
    self.use_batchnorm = use_batchnorm
    self.use_dropout = dropout > 0
    self.reg = reg
    self.num_layers = 1 + len(hidden_dims)
    self.dtype = dtype
    self.params = {}

    # ===== #
    # YOUR CODE HERE:
    # Initialize all parameters of the network in the self.params
    dictionary.
    # The weights and biases of layer 1 are W1 and b1; and in general the
    # weights and biases of layer i are Wi and bi. The
    # biases are initialized to zero and the weights are initialized
    # so that each parameter has mean 0 and standard deviation
    weight_scale.
    # ===== #
    for i in np.arange(self.num_layers):

```

```

        if(i == 0):
            self.params['W' + str(i+1)] = np.random.normal(0, weight_scale,
↪(input_dim, hidden_dims[i]))
            self.params['b' + str(i+1)] = np.zeros(hidden_dims[i])
        elif(i == self.num_layers - 1):
            self.params['W' + str(i+1)] = np.random.normal(0, weight_scale,
↪(hidden_dims[i-1], num_classes))
            self.params['b' + str(i+1)] = np.zeros(num_classes)
        else:
            self.params['W' + str(i+1)] = np.random.normal(0, weight_scale,
↪(hidden_dims[i-1], hidden_dims[i]))
            self.params['b' + str(i+1)] = np.zeros(hidden_dims[i])

    #pass

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    # When using dropout we need to pass a dropout_param dictionary to each
    # dropout layer so that the layer knows the dropout probability and the
↪mode
    # (train / test). You can pass the same dropout_param to each dropout
↪layer.
    self.dropout_param = {}
    if self.use_dropout:
        self.dropout_param = {'mode': 'train', 'p': dropout}
        if seed is not None:
            self.dropout_param['seed'] = seed

    # With batch normalization we need to keep track of running means and
    # variances, so we need to pass a special bn_param object to each batch
    # normalization layer. You should pass self.bn_params[0] to the forward
↪pass
    # of the first batch normalization layer, self.bn_params[1] to the
↪forward
    # pass of the second batch normalization layer, etc.
    self.bn_params = []
    if self.use_batchnorm:
        self.bn_params = [{'mode': 'train'} for i in np.arange(self.
↪num_layers - 1)]

    # Cast all parameters to the correct datatype
    for k, v in self.params.items():
        self.params[k] = v.astype(dtype)

```

```

def loss(self, X, y=None):
    """
    Compute loss and gradient for the fully-connected net.

    Input / output: Same as TwoLayerNet above.
    """
    X = X.astype(self.dtype)
    mode = 'test' if y is None else 'train'

    # Set train/test mode for batchnorm params and dropout param since they
    # behave differently during training and testing.
    if self.dropout_param is not None:
        self.dropout_param['mode'] = mode
    if self.use_batchnorm:
        for bn_param in self.bn_params:
            bn_param[mode] = mode

    scores = None

    # ===== #
    # YOUR CODE HERE:
    # Implement the forward pass of the FC net and store the output
    # scores as the variable "scores".
    # ===== #
    H = []
    H_cache = []
    for i in range(self.num_layers):
        H_app = None
        H_cache_app = None
        if(i==0):
            H_app, H_cache_app = affine_relu_forward(X, self.params['W' +
↪str(i+1)], self.params['b' + str(i+1)])
            H.append(H_app)
            H_cache.append(H_cache_app)
        elif(i == self.num_layers - 1):
            scores, H_cache_app = affine_forward(H[i-1], self.params['W' +
↪str(i+1)], self.params['b' + str(i+1)])
            H_cache.append(H_cache_app)
        else:
            H_app, H_cache_app = affine_relu_forward(H[i-1], self.
↪params['W' + str(i+1)], self.params['b' + str(i+1)])
            H.append(H_app)
            H_cache.append(H_cache_app)

    #pass

    # ===== #

```

```

# END YOUR CODE HERE
# ===== #

# If test mode return early
if mode == 'test':
    return scores

loss, grads = 0.0, {}
# ===== #
# YOUR CODE HERE:
# Implement the backwards pass of the FC net and store the gradients
# in the grads dict, so that grads[k] is the gradient of self.
↪params[k]
# Be sure your L2 regularization includes a 0.5 factor.
# ===== #

#pass
loss, dhidden = softmax_loss(scores, y)
for i in range(self.num_layers, 0, -1):
    loss += 0.5*self.reg*np.sum(self.params['W{}'.format(i)]*self.
↪params['W{}'.format(i)])
    if i == self.num_layers:
        dH1, dW, db = affine_backward(dhidden, H_cache[i-1])
        grads['W{}'.format(i)] = dW + self.reg*self.params['W{}'.format(i)]
        grads['b{}'.format(i)] = db
    else:
        dH1, dW, db = affine_relu_backward(dH1, H_cache[i-1])
        grads['W{}'.format(i)] = dW + self.reg*self.params['W{}'.format(i)]
        grads['b{}'.format(i)] = db
# loss, dhidden = softmax_loss(scores, y)
# for i in range(self.num_layers, 0, -1):
#     loss += 0.5 * self.reg*np.sum(self.params['W{}'.format(i)]*self.
↪params['W{}'.format(i)])
#     if i == self.num_layers:
#         dFC1, dW, db = affine_backward(dhidden, FC_cache[i-1])
#         grads['W{}'.format(i)] = dW + self.reg*self.params['W{}'.format(i)].
↪format(i)]
#         grads['b{}'.format(i)] = db
#     else:
#         dFC1, dW, db = affine_relu_backward(dFC1, FC_cache[i-1])
#         grads['W{}'.format(i)] = dW + self.reg*self.params['W{}'.format(i)].
↪format(i)]
#         grads['b{}'.format(i)] = db
# ===== #
# END YOUR CODE HERE
# ===== #
return loss, grads

```